

Advanced-Level

Architecture & Design

1. How would you design a highly available and fault-tolerant web application on AWS?

- Use **Multi-AZ VPC** with public (LB) and private (app/DB) subnets.
- Put an **Application Load Balancer (ALB)** in front to distribute traffic.
- Run app servers on **EC2 Auto Scaling, ECS, or EKS** across **multiple AZs**.
- Use **RDS/Aurora Multi-AZ** for a highly available database.
- Store static files in **S3**, optionally serve via **CloudFront**.
- Add **ElastiCache (Redis/Memcached)** for caching.
- Use **SQS/SNS** for decoupling background tasks.
- Secure with **WAF, Security Groups, IAM, and KMS**.
- Monitor via **CloudWatch, X-Ray, CloudTrail**.
- Use **Route 53** for DNS failover and DR.

Multi-AZ, Auto-scaling, Load Balancing, Decoupling, Managed Services, and Self-healing infrastructure.

1. Network Layer (VPC Design)

- Create a **VPC** with at least:
 - **2 public subnets** (across 2 Availability Zones)
 - **2 private subnets** (across 2 Availability Zones)
- Put **Load Balancers in public subnets**, application servers + database in private subnets.
- Use **NAT Gateways** in each AZ for outbound traffic.

Benefit: Multi-AZ networking ensures no single point of failure.

2. Load Balancing Layer

Use **Application Load Balancer (ALB)**:

- Distribute traffic across multiple EC2 instances or containers
- Health checks remove unhealthy instances automatically
- Supports path-based routing for microservices

Benefit: High availability + intelligent routing.

3. Compute Layer

Option A: EC2 Auto Scaling Group

- Minimum 2 instances across 2 AZs
- Auto scale based on CPU, requests, memory (via CloudWatch agent)
- Launch templates for consistency

Option B (preferred): Containerized Compute

With ECS (Fargate):

- Serverless containers across multiple AZs
- Service Auto Scaling

With EKS (Kubernetes):

- Multi-AZ node groups
- Self-healing workloads

Benefit: Self-healing compute with multi-AZ redundancy.

4. Database Layer

Option A: Amazon RDS (MySQL/PostgreSQL/MariaDB)

- Multi-AZ deployment with synchronous replication
- Automatic failover to standby node
- Backups and snapshots

Option B: Amazon Aurora

- Storage spans 3 AZs automatically
- Up to 15 read replicas
- Faster failover (~30 seconds or less)

Option C: DynamoDB

- Fully managed, multi-AZ by default
- Point-in-time recovery + global tables

Benefit: Zero-admin, multi-AZ, highly available DB.

5. Caching Layer

Amazon ElastiCache (Redis or Memcached)

- Multi-AZ with automatic failover (Redis)

- Reduces read load on the database
- Speeds up response times

6. Storage Layer

Amazon S3

- Store static assets (images, PDFs, static websites)
- Enable **S3 versioning**, **S3 replication**, **lifecycle policies**
- Optional **CloudFront CDN** to serve global traffic

Benefit: 99.99999999% durability + low latency.

7. Decoupling Layer

Use managed messaging to avoid tightly coupled systems:

Amazon SQS

- Queue for asynchronous processing

Amazon SNS

- Pub/Sub messaging to notify services

Amazon EventBridge

- Event-driven integration between microservices

Benefit: Fault isolation and resilience.

8. Monitoring, Logging & Alerts

Use:

- **Amazon CloudWatch** (metrics, alarms, dashboards)
- **AWS X-Ray** (distributed tracing)
- **CloudTrail** (API governance)
- **ELB access logs & S3 logs**
- **OpenSearch** for log indexing (optional)

Benefit: Faster diagnosis and proactive scaling.

9. Security & IAM Controls

- **WAF + Shield** for Layer 7/Layer 3 DDoS protection
- **Security Groups** for instance-level protection
- **NACLs** for subnet control

- **Secrets Manager** for DB passwords + API keys
- **KMS** for encryption at rest

Benefit: Strong, layered security.

10. Disaster Recovery Design

Depending on business requirement:

DR Strategy Options:

Strategy	RTO	RPO	Uses
Backup & Restore	Hours	Hours	S3 + Snapshots
Pilot Light	Minutes	Minutes	Minimal resources in secondary region
Warm Standby	Seconds–Minutes	Seconds	Scaled-down full stack
Multi-Region Active/Active	Near zero	Near zero	Global apps w/ Route 53 Geolocation

Tools for DR:

- Route 53 failover routing
- Cross-region replication (S3, DynamoDB)
- RDS cross-region read replicas
- CloudFront global edge locations

Final Architecture Summary

Your **Highly Available & Fault-Tolerant AWS Web App** includes:

- Multi-AZ VPC design
- **Load Balancer** for traffic distribution
- **EC2 Auto Scaling, ECS, or EKS** compute
- **RDS/Aurora** with Multi-AZ
- **ElastiCache** for caching
- **S3 + CloudFront** for static content
- **SQS/SNS** for decoupling
- **WAF + Shield + IAM** for security
- **CloudWatch, X-Ray, CloudTrail** for monitoring
- **Route 53** for global DNS and failover
- **DR with multi-region replication**

2. How do you achieve multi-region deployment in AWS?

1. **Deploy full application stack in multiple AWS regions**
(VPC, ALB, ECS/EC2/EKS, RDS/Aurora, S3, etc.)
2. **Use Route 53 for global traffic routing**

- Latency-based routing
 - Geolocation routing
 - Failover routing for DR
 - Route 53 directs users to the nearest healthy region.
- 3. Replicate data across regions**
- **Aurora Global Database** or RDS cross-region replicas
 - **DynamoDB Global Tables** (fully multi-master)
 - **S3 Cross-Region Replication** (CRR)
 - **ElastiCache Global Datastore** (for Redis)
- 4. Sync application configuration and secrets**
- Parameter Store / Secrets Manager multi-region replication
 - CI/CD pipelines deploying to all regions
- 5. Decouple services with global messaging**
- SQS with Lambda fan-out
 - SNS cross-region topics
 - EventBridge global event bus (if needed)
- 6. Use CloudFront CDN for edge distribution**
Delivers static assets from the closest edge location.
- 7. Implement health checks and failover**
- Route 53 health checks
 - Automatic failover from primary to secondary region

Summary

Use **multi-region deployments**, **Route 53 global routing**, and **cross-region data replication** to ensure your application continues running even if an entire AWS region fails.

3. How do you handle disaster recovery in AWS?

AWS provides several DR strategies based on cost, RTO (recovery time), and RPO (data loss tolerance). The key methods are:

1. Backup & Restore (Lowest cost)

- Store backups in **S3** or **Glacier**
- Use **RDS snapshots**, **EBS snapshots**, **S3 versioning**
- Restore resources in another region when needed

RTO: Hours | RPO: Hours

2. Pilot Light

- Keep **core services** (database, minimal servers) running in a second region
- Spin up the full environment during DR
- Use **replicated databases**, **AMI copies**, **IaC** to launch quickly

RTO: Minutes–Hours | RPO: Minutes

3. Warm Standby

- Scaled-down version of the full app running in another region
- Scale up to full capacity during DR
- Use **automated failover** with Route 53

RTO: Minutes | RPO: Seconds–Minutes

4. Multi-Region Active/Active (Highest availability)

- Application runs in **multiple regions simultaneously**
- Use:
 - **Route 53 latency routing / failover routing**
 - **Aurora Global Database / DynamoDB Global Tables**
 - **S3 Cross-Region Replication**
- Traffic automatically shifts if a region fails

RTO: Near zero | RPO: Near zero

Core AWS Tools Used for DR

- **Route 53 health checks & failover**
- **Cross-Region Replication** (S3, DynamoDB, Aurora, RDS replicas)
- **CloudFormation/Terraform** for rapid rebuild
- **CloudWatch alarms** for incident detection
- **AWS Backup** for centralized backup management
- **SNS** for DR notifications

Short Summary

“Disaster recovery in AWS is achieved using multi-region architecture, automated backups, cross-region replication, and Route 53 failover. Depending on business needs, AWS supports Backup & Restore, Pilot Light, Warm Standby, and Active/Active strategies to meet specific RTO and RPO requirements.”

4. How do you optimize cost in AWS?

1. Right-size Resources

- Identify underutilized EC2, RDS, EBS using **CloudWatch, Cost Explorer, Compute Optimizer**
- Choose correct instance sizes, families, and storage types

- Stop or terminate idle resources

2. Use Auto Scaling

- Scale **up and down automatically** based on load
- Avoid paying for unused capacity

3. Use Savings Plans & Reserved Instances

- **Compute Savings Plans** for EC2, Fargate, Lambda
- **Reserved Instances (RIs)** for steady workloads (RDS, EC2, ElastiCache)
- Up to **70% cost reduction**

4. Use Spot Instances

- Use **EC2 Spot Instances** for batch, CI/CD, analytics, or non-critical workloads
- Up to **90% cheaper** than On-Demand

5. Optimize Storage Costs

- Use **S3 lifecycle policies** → move data to Standard-IA, Glacier, Deep Archive
- Delete unused EBS volumes & snapshots
- Use **EFS IA** for infrequently accessed data

6. Choose Cost-Effective Databases

- Use **Aurora Serverless** or DynamoDB for variable workloads
- Use **RDS storage auto-scaling**
- Turn on **RDS pause/resume** for dev databases

7. Use Serverless Where Possible

- Lambda, API Gateway, Fargate
- Pay only for usage, not provisioned capacity

8. Use CloudFront for Caching

- Reduces load on origin servers
- Speeds up delivery + reduces data transfer cost

9. Turn Off Non-Prod Resources

- Use automation to shut down dev/test environments after business hours

10. Monitor & Enforce Budgets

- Use **AWS Budgets** + alerts
- Use **Cost Explorer** to track spend

- Tag resources for accountability

Short Summary

“To optimize cost in AWS, I right-size resources, use auto-scaling, adopt Reserved Instances or Savings Plans, leverage Spot Instances, optimize storage with lifecycle rules, move workloads to serverless where possible, and use monitoring tools like Cost Explorer and AWS Budgets to track and control spending.”

DevOps & Automation

1.What is Infrastructure as Code (IaC)?

Infrastructure as Code (IaC) is the practice of **managing and provisioning infrastructure (servers, networks, databases, load balancers, etc.) using machine-readable configuration files instead of manual processes.**

In simple terms:

IaC lets you **write code to create and manage your cloud infrastructure**, just like you write code to build an application.

Key Benefits

- **Consistency:** No manual errors; same setup every time
- **Automation:** Faster deployments and easy rollbacks
- **Scalability:** Quickly spin up or tear down environments
- **Version Control:** Track changes using Git
- **Cost Efficiency:** Automatically remove unused resources

Common IaC Tools

- **Terraform**
- **AWS CloudFormation**
- **Pulumi**
- **Ansible**
- **Chef / Puppet**

In short

“Infrastructure as Code is the practice of defining and managing infrastructure using code instead of manual setup. It makes deployments consistent, automated, repeatable, and version-controlled.”

2. Compare AWS CloudFormation and Terraform.

1. Ecosystem & Vendor Lock-in

CloudFormation

- **AWS-native** IaC service.
- Best aligned with AWS features, often gets support for new AWS services first.
- **Vendor-locked**: cannot manage resources outside AWS (except limited 3rd-party integrations via the AWS CloudFormation Registry).

Terraform

- **Multi-cloud and multi-provider**: AWS, Azure, GCP, GitHub, Kubernetes, Datadog, etc.
- Reduces vendor lock-in, supports hybrid cloud and large tool ecosystems.

Winner: *Terraform for flexibility; CloudFormation if you're fully committed to AWS.*

2. Language / Syntax

CloudFormation

- Uses **YAML or JSON**, which can become verbose.
- AWS added **CloudFormation Modules** and **Serverless Application Model (SAM)** to offset complexity, but still less ergonomic.

Terraform

- Uses **HCL (HashiCorp Configuration Language)**.
- More concise, readable, and modular than raw YAML/JSON.

Winner: *Terraform (HCL is easier and more expressive).*

3. State Management

CloudFormation

- **No external state file**; state is fully managed by AWS.
- Safer in terms of not losing/corrupting state.
- Harder to reference resources created outside CloudFormation.

Terraform

- Maintains a **state file** (local or remote).
- Supports **Terraform Cloud**, S3, and other backends for remote locking.

- Very flexible, but mismanaging state can cause issues.

Winner: *CloudFormation* for simplicity; *Terraform* for flexibility.

4. Modularity & Reusability

CloudFormation

- Supports **Nested Stacks, Modules, Macros, and Transforms**.
- Still more limited and verbose.

Terraform

- Strong **module system**, easily shareable via the Terraform Registry.
- Built-in variable and output systems are clean and powerful.

Winner: *Terraform*.

5. Plan & Apply Workflow

CloudFormation

- “Change sets” allow you to preview changes, but the UX is clunkier.

Terraform

- `terraform plan` clearly shows detailed changes.
- `terraform apply` then performs the deployment.
- Excellent clarity and developer experience.

Winner: *Terraform*.

6. Drift Detection

CloudFormation

- Built-in **drift detection** capability.

Terraform

- Supports drift detection through `terraform plan` and by refreshing state, but it’s less explicit.

Winner: *CloudFormation*.

7. Community & Ecosystem

CloudFormation

- Backed by AWS, with strong documentation, but more narrow ecosystem.

Terraform

- Massive open-source community.
- Thousands of modules and providers.
- Lots of tooling and integration options.

Winner: *Terraform*.

[8. Cost](#)

CloudFormation

- **Free** (you only pay for AWS resources you create).

Terraform

- **Open-source Core is free.**
- Terraform Cloud/Enterprise costs extra (optional).

Winner: *CloudFormation (if you count every dollar), but Terraform's free OSS version is enough for most teams.*

When to Choose What

[Choose CloudFormation if:](#)

- You are **100% AWS-only**.
- You value **built-in state management**.
- You want **tighter AWS integration**.
- You prefer **fully AWS-managed** solutions.

[Choose Terraform if:](#)

- You operate in **multi-cloud or hybrid** environments.
- You want **cleaner syntax**, better modularity, and a great workflow.
- You rely on many external SaaS providers.
- You want to standardize IaC across different platforms

How does AWS CodePipeline work?

Summary Table

Feature	CloudFormation	Terraform
Cloud support	AWS only	Multi-cloud
Language	YAML/JSON	HCL
State management	Managed by AWS	External, flexible
Drift detection	Built-in	Indirect
Modularity	Decent, limited	Excellent
Plan/apply UX	Good	Excellent
Community	AWS-centric	Huge OSS ecosystem
Cost	Free	Mostly free

3. How does AWS CodePipeline work?

AWS CodePipeline is a **fully managed continuous integration and continuous delivery (CI/CD) service** that automates your software release process. It helps you move code from commit → build → test → deploy in a repeatable, reliable way.

How AWS CodePipeline Works

1. Pipelines are defined as a series of stages

A pipeline is made up of stages. Common stages include:

1. **Source** – detect changes in your code repository
2. **Build** – compile, run tests, create artifacts
3. **Test** – optional integration or automated tests
4. **Deploy** – publish the application to environments (e.g., EC2, ECS, Lambda, S3)

Each stage contains **actions**, which do the actual work (like pulling source, running builds, or deploying).

2. Source Stage: Triggering the Pipeline

CodePipeline integrates with sources such as:

- AWS CodeCommit
- GitHub / GitHub Enterprise
- Bitbucket
- Amazon S3

When a change is detected, the pipeline **automatically triggers**.

3. Build Stage: Using AWS CodeBuild (or others)

The build stage commonly uses:

- **AWS CodeBuild** (fully managed build service)
- Jenkins
- TeamCity
- Custom build actions

CodeBuild compiles code, runs tests, and creates deployment **artifacts** stored in Amazon S3.

4. Test Stage (optional)

You can include automated tests such as:

- Integration tests
- UI tests
- Load tests

These can run using:

- CodeBuild
- Lambda functions
- Third-party systems

5. Deploy Stage

CodePipeline deploys artifacts to environments such as:

- **AWS Elastic Beanstalk**
- **Amazon ECS / EKS**
- **AWS Lambda**
- **Amazon EC2 (via CodeDeploy)**
- **Amazon S3** (for static sites)
- **CloudFormation** for infrastructure deployments

Deployment strategies like **blue/green** or **rolling updates** can be configured.

6. Pipeline Orchestration & Automation

CodePipeline orchestrates stages, passing artifacts and metadata between them. It includes features like:

- **Manual approval actions** (e.g., approve before deploying to production)
- **Parallel actions** (multiple builds or tests at once)
- **Retry handling**
- **Event-driven hooks** (via Amazon EventBridge)

7. Pipeline as Code

You can define your pipeline using:

- AWS Management Console (visual designer)
- CloudFormation
- AWS CDK
- Terraform

This makes your CI/CD configuration version-controlled and reproducible.

High-Level Flow

1. **Developer pushes code**
2. **Source stage triggers**
3. **CodeBuild compiles & tests**
4. **Artifacts generated**
5. **Automated tests run**
6. **Deployment to environment**
7. **Manual approval (optional)**
8. **Production deployment**

Example Use Case

When you push code to GitHub:

1. GitHub notifies CodePipeline
2. CodePipeline starts
3. CodeBuild runs tests → uploads artifact
4. CodeDeploy updates EC2 instances
5. The application is live

4. How do you integrate CI/CD with AWS services?

Integrating **CI/CD with AWS services** typically means using AWS's native developer tools (or mixing them with third-party tools) to automate the process of building, testing, and deploying applications. Below is a clear end-to-end explanation of *how to integrate CI/CD on AWS*, including common architectures and the services involved.

Core AWS Services Used for CI/CD

AWS provides a suite of tools that integrate naturally:

Stage	AWS Service	Purpose	🔗
Source	CodeCommit / GitHub / Bitbucket / S3	Stores source code	
Build	CodeBuild	Compiles code, runs tests	
Pipeline Orchestration	CodePipeline	Automates CI/CD workflow	
Deploy	CodeDeploy, Elastic Beanstalk, ECS, EKS, Lambda, CloudFormation	Deploys artifacts	
Artifact Storage	S3	Stores build outputs	
Monitoring	CloudWatch / CloudTrail	Logs and audit trails	

How CI/CD Integration Works on AWS

1. Set Up Your Source Repository

Your repo can be:

- **AWS CodeCommit**
- **GitHub / GitHub Enterprise**
- **Bitbucket**
- **Amazon S3 (for zipped artifacts)**

When code is pushed, an event triggers the pipeline automatically.

2. Build Stage with AWS CodeBuild

CodeBuild is a fully managed build environment.

You define a **buildspec.yml** in your repo, for example:

```
version: 0.2
phases:
  install:
    commands:
      - npm install
  build:
    commands:
      - npm run build
artifacts:
  files:
    - '**/*'
```

CodeBuild outputs test results and build artifacts to **S3**.

3. Create a CI/CD Pipeline Using AWS CodePipeline

CodePipeline is where the full workflow is orchestrated.

A typical pipeline includes:

Source → Build → Test → Deploy

You can add:

- **Manual approval** before deploying to production
- **Parallel actions** for multiple environments
- **Automatic rollback** if deployment fails

4. Deployment Options

Based on your architecture, you integrate CodePipeline with a deployment service:

a. EC2 or On-Premise Servers → CodeDeploy

Supports:

- Rolling updates
- Blue/green deployments
- Canary deployments

b. Containerized apps → ECS or EKS

- Deploy via **CodePipeline** → **CodeBuild** → **ECS Blue/Green (via CodeDeploy)**
- For EKS, usually CodePipeline triggers a `kubectl apply` via CodeBuild.

c. Serverless apps → Lambda

Deploy via:

- **CodePipeline + Lambda function versioning**
- **SAM (Serverless Application Model)**
- **CloudFormation** for IaC deployments

d. Web/frontend apps → S3 + CloudFront

- Build in CodeBuild
- Upload to S3
- Invalidate CloudFront cache in deployment action

5. Infrastructure as Code (IaC) Integration

Best practice: Deploy your **infrastructure** also through the CI/CD pipeline.

Options:

- **AWS CloudFormation**
- **AWS CDK**
- **Terraform**

This ensures infra + application deploy together in a consistent way.

Example: Common AWS CI/CD Architecture

For an ECS (Docker) application:

1. Developer pushes code to GitHub
2. GitHub triggers CodePipeline
3. **CodeBuild:**
 - Runs tests
 - Builds Docker image
 - Pushes image to ECR
4. CodeDeploy performs a **blue/green** ECS deployment
5. Load balancer shifts traffic to new tasks

Third-Party CI/CD Tool Integration

AWS integrates easily with non-AWS CI/CD systems:

GitHub Actions → CodeDeploy/ECS/Lambda

Use AWS GitHub Actions for:

- aws-actions/configure-aws-credentials
- aws-actions/amazon-ecr-login

GitLab CI → AWS

Use GitLab runners pushing artifacts or Docker images to ECR, and deploy with CloudFormation/CodeDeploy.

Jenkins → AWS

Use AWS plugin to trigger CodeBuild/CodeDeploy.

Automated Testing Integration

You can add automated tests:

- **Unit tests** → CodeBuild
- **Integration tests** → Lambda or CodeBuild
- **End-to-end tests** → Selenium or Cypress running in CodeBuild container

Observability & Monitoring

AWS provides:

- **CloudWatch Logs** (build/test logs)
- **CloudTrail** (audit pipeline activity)
- **CloudWatch Alarms** (trigger rollback or notifications)
- **SNS** (pipeline notifications)

Want a Deployment Example?

I can provide:

❖ A full CI/CD example for:

- **ECS**
- **Lambda**
- **EC2/CodeDeploy**
- **Static website (S3/CloudFront)**
- **Serverless stack (SAM/CDK)**

or a full **CodePipeline YAML/CloudFormation/CDK template**.

Security & Compliance

1. How do you secure an S3 bucket to make it private but still accessible by certain users?

To make an Amazon S3 bucket **private** while still allowing **specific users to access it**, you typically combine **Block Public Access**, **bucket policies**, and **IAM user/role permissions**.

1. Block All Public Access (Make Bucket Private)

In the S3 console:

1. Go to **Bucket → Permissions → Block public access**.
2. Turn on **all** of these:
 - Block public ACLs
 - Block public bucket policies
 - Block new public ACLs
 - Block new public policies

This ensures **nobody** gets public/anonymous access.

↙2. Remove Any Existing Public ACLs

- S3 ACLs should normally be disabled or set to **Bucket owner full control**.
- Make sure no ACL grants access to *Everyone (public)*.

. Grant Access to Specific IAM Users or Roles

we have two main options:

Option A — Grant Access Using IAM Policies (Recommended)

Attach a policy *to the IAM user or role* that needs access:

Example: Allow read-only access to one bucket

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": [  
        "s3:GetObject",  
        "s3>ListBucket"  
      ],  
      "Resource": [  
        "arn:aws:s3:::my-private-bucket",  
        "arn:aws:s3:::my-private-bucket/*"  
      ]  
    }  
  ]  
}
```

Example: Allow read/write access

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": "s3:*",  
      "Resource": [  
        "arn:aws:s3:::my-private-bucket",  
        "arn:aws:s3:::my-private-bucket/*"  
      ]  
    }  
  ]  
}
```

This is the most maintainable method.

Option B — Use a Bucket Policy to Allow Only Certain Users

If you want access controlled *at the bucket level*, you can add this bucket policy:

Example: Allow access only to one specific IAM user or role

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Principal": "arn:aws:iam::123456789012:User/jdoe",  
      "Action": "s3:GetObject",  
      "Resource": "arn:aws:s3:::my-private-bucket/*"  
    }  
  ]  
}
```

```
"Effect": "Allow",
"Principal": {
    "AWS": "arn:aws:iam::123456789012:user/Bob"
},
>Action": [
    "s3:GetObject",
    "s3>ListBucket"
],
"Resource": [
    "arn:aws:s3:::my-private-bucket",
    "arn:aws:s3:::my-private-bucket/*"
]
}
]
```

This works even with Block Public Access on because the user is *not* public.

When to Use IAM Policies vs Bucket Policies

Goal	Best Method
Manage access per user/role	IAM policy
Restrict access from inside the bucket itself	Bucket policy
Cross-account access	Bucket policy (+ possibly IAM roles)

4. For Temporary or Signed Access — Use Pre-Signed URLs

If someone needs temporary access without being an AWS user:

```
aws s3 presign s3://my-private-bucket/file.pdf --expires-in 3600
```

short

- Block public access
- Remove public ACLs
- Use IAM policies or bucket policies to specify allowed users
- Use pre-signed URLs for temporary access

2.What are AWS Organizations and Service Control Policies (SCPs)?

AWS Organizations

AWS Organizations is a service that lets you centrally manage and govern multiple AWS accounts. It is used for:

What it helps you do

- **Create and manage multiple AWS accounts** within a unified structure
- **Group accounts** using Organizational Units (OUs)
- **Apply policies** across accounts (e.g., SCPs)
- **Enable consolidated billing** (all accounts roll up to one payer account)
- **Improve security** by isolating workloads into separate accounts
- **Delegate administration** of certain AWS services (e.g., IAM Identity Center, CloudTrail)

Why it's useful

- Strong isolation between accounts
- Least privilege across environments (e.g., dev/stage/prod)
- Centralized governance and guardrails
- Reduced blast radius if one account is compromised

Service Control Policies (SCPs)

SCPs are policies that set the maximum permissions allowed in AWS accounts or Organizational Units (OUs) within an AWS Organization.

Think of SCPs as **permission guardrails**.

Key Concepts

- SCPs **do not grant permissions**.
- They **limit** what IAM users, roles, and even the root user in an account can do.
- If an action is **not allowed by an SCP**, no IAM policy can enable it.
- Applied at the **OU or account level**.
- Work in combination with IAM policies → both must allow an action.

Example use cases

- Enforce that accounts **cannot disable CloudTrail**
- Block use of expensive or high-risk services
- Restrict regions (e.g., "Only allow us-east-1 and us-west-2")
- Prevent creation of IAM users (force IAM Identity Center usage)
- Require encryption for resources

Example SCP: Deny use of all services except S3 and EC2

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Deny",  
            "NotAction": [  
                "s3:*",  
                "ec2:*"  
            ],  
            "Resource": "*"  
        }  
    ]  
}
```

This means:

- Only S3 and EC2 actions will be allowed
- Even root user cannot use anything else

Example SCP: Deny operations outside certain regions

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Deny",  
            "Action": "*",  
            "Resource": "*",  
            "Condition": {  
                "StringNotEquals": {  
                    "aws:RequestedRegion": ["us-east-1", "us-west-2"]  
                }  
            }  
        }  
    ]  
}
```

Mental Model

- IAM policies = permissions you *grant*
- SCPs = permissions you *can't exceed*

Final permission = IAM ALLOW AND SCP ALLOW
→ If either denies, access is denied.

Summary

AWS Organizations

- Structure and manage multiple accounts
- Group accounts in OUs
- Centralized billing and governance

Service Control Policies (SCPs)

- Account-level guardrails
- Restrict actions even for root
- Do not grant permissions — only limit them

3. How do you manage secrets and credentials in AWS?

1. AWS Secrets Manager (Best for passwords, API keys, database creds)

AWS Secrets Manager is the primary service for storing and retrieving sensitive credentials.

What it provides

- Encrypted storage (via KMS)
- Automatic rotation (built-in for RDS, Redshift, DocumentDB; custom Lambda for others)
- Fine-grained IAM access control
- CloudTrail auditing of secret use
- Versioning and staged rotation

How to retrieve a secret (example)

```
import boto3

client = boto3.client("secretsmanager")
secret = client.get_secret_value(SecretId="my-db-secret") ["SecretString"]
```

When to use it:

- Database credentials
- API keys
- OAuth tokens
- Third-party service credentials

2. AWS Systems Manager Parameter Store

Parameter Store supports two types of parameters:

- **Standard** (free)
- **Advanced** (paid, with extra features)

Features

- Encrypted parameters using KMS (SecureString)
- Hierarchical parameter paths (e.g., /prod/db/password)
- IAM access control
- Good for config + secrets (not as rich as Secrets Manager)

Example retrieving secure parameter:

```
aws ssm get-parameter --name "/prod/db/password" --with-decryption
```

When to use it:

- App configuration
- Secrets with less strict rotation needs
- Free/cheaper alternative to Secrets Manager

3. IAM Roles & Temporary Credentials (Best for AWS service access)

Instead of storing AWS API keys, you **never use long-term credentials** for workloads.

Use IAM roles:

- EC2 Instance Roles
- Lambda Execution Roles
- ECS Task Roles
- EKS IRSA Roles (web identity)

These provide **temporary credentials** automatically rotated by AWS STS.

Why:

- No hard-coded access keys
- Auto-rotation every few minutes
- Scoped least-privilege permissions

4. AWS KMS (Encryption for secrets)

KMS is not a secret store, but it is used to:

- Encrypt secrets stored in S3 or Parameter Store
- Create Customer-Managed Keys (CMKs) for compliance
- Control access using IAM + key policies

Example:

- Encrypting config files
- Encrypting values before storing in DynamoDB or S3

5. AWS Cognito for user authentication

For external-facing apps (mobile/web), **Cognito handles user credentials** so you don't store them yourself.

- User pools for identity
- Federated login (Google, Apple, etc.)
- Secure token handling (JWTs)
- MFA and password policies

6. Best Practices (What you should *always* do)

✓ Never store secrets in code, GitHub, or environment files

Use Secret Manager or Parameter Store.

✓ Use IAM roles instead of access keys

Eliminate long-term AWS keys entirely.

✓ Enable automatic rotation for database credentials and secrets where possible.

✓ Restrict access using IAM

Use fine-grained policies:

- <app-role> can only read <app-secret>

✓ Audit everything with CloudTrail

You should always know:

- Who retrieved a secret
- When
- From where

✓ Use resource policies for cross-account access

Secrets Manager supports resource-based access control.

Which Service Should You Use? (Quick Comparison)

Use case	Best choice
Database username/password	Secrets Manager
API keys or OAuth tokens	Secrets Manager
General config values	Parameter Store
Cost-sensitive environments	Parameter Store (Standard)
Workload-to-AWS authentication	IAM Roles
Encryption of your own stored secrets (S3, DynamoDB)	KMS
External user login for an app	Cognito

Performance & Scaling

1. How do you scale databases in AWS?

Scaling databases in AWS can be done **vertically**, **horizontally**, or through **managed distributed services**. The right strategy depends on your workload (read-heavy, write-heavy, analytics, global traffic, etc.)

1. Vertical Scaling (Scale Up/Down)

Increase the instance size/resources of your database server.

Works with:

- **Amazon RDS** (MySQL, Postgres, MariaDB, Oracle, SQL Server)
- **Amazon Aurora**
- **EC2 self-managed databases**

What it means

- Increase CPU, RAM, storage type (gp3 → io2), or IOPS.
- Easy, but limited: you're bound by the largest instance size.

When to use

- When performance issues come from insufficient compute or memory.
- Quick fix for predictable workloads.

2. Horizontal Scaling (Scale Out)

Add more database nodes/replicas to distribute load.

A. Read Scaling – Add Read Replicas

Offload read queries from the primary database.

Works with:

- RDS (MySQL, PostgreSQL, MariaDB)
- Aurora (best replica architecture)
- DynamoDB (auto-scaled reads)

Benefits

- Reduce load on primary
- Geographic distribution (global read replicas)
- Failover capabilities

B. Sharding / Partitioning

Split data across multiple independent databases.

Approaches:

- **Application-level sharding** (your code chooses the shard)
- **Aurora global database** (multi-region low-latency reads)
- **DynamoDB partitioning** (automatic)

Use cases

- Very large datasets (TB–PB)
- High-write workloads
- Multi-tenant applications (each tenant on a shard)

3. Managed Distributed Databases

AWS offers databases that **scale automatically** with minimal management.

A. Amazon Aurora (MySQL/Postgres-compatible)

- Up to 15 read replicas
- Auto-scaling storage (10GB → 128TB)
- Serverless v2 supports near-instant autoscaling

Best for: general-purpose OLTP, high availability, auto-scaling compute.

B. Amazon DynamoDB (Fully serverless NoSQL)

- Infinite read/write scaling
- Auto-scaling throughput
- Global tables for multi-region replication
- Millisecond performance even at millions of TPS

Best for: massive scale, high traffic, key-value workloads.

C. Amazon Redshift / Redshift Serverless

- Scales compute and storage independently
- Distributed MPP architecture
- Great for analytics, BI

Best for: OLAP workloads and large-scale analytics.

4. Caching Layers (Reduce Database Load)

Caching is the most efficient way to scale read-heavy workloads.

A. Amazon ElastiCache (Redis / Memcached)

- In-memory caching for low latency (microseconds)
- Helps offload repetitive queries
- Can store sessions, counters, partial query results

B. DynamoDB Accelerator (DAX)

- 10x faster DynamoDB reads
- Microsecond latency
- Drop-in caching layer

5. Global Scaling

A. Aurora Global Database

- Low-latency global reads
- Fast disaster recovery

B. DynamoDB Global Tables

- Multi-region multi-master replication
- Active-active traffic globally

C. Amazon RDS Cross-Region Replicas

- Read replicas across regions

6. Event-Driven & Async Patterns (Database Offloading)

Reduce write pressure by moving expensive work to asynchronous systems.

Tools:

- **Amazon SQS + Lambda**
- **Amazon Kinesis / MSK (Kafka)**
- **EventBridge**

Use cases:

- Write-intensive workloads
- Analytics pipelines
- Background processing

This pattern helps databases scale by smoothing out spikes.

7. Storage Scaling

Aurora

- Auto scales storage

RDS

- Increase disk or switch to high-performance storage types (io2)

DynamoDB

- Storage scales automatically

8. Auto-Failover and High Availability

Not directly “scaling,” but required for scale.

Examples:

- Multi-AZ deployments (RDS, Aurora)
- Aurora automatic failover (fast < 30s)
- DynamoDB multi-region HA

High availability allows the system to scale without interruptions.

2. What are caching strategies in AWS (CloudFront, ElastiCache, etc.)?

🎯 Choosing the Right AWS Caching Strategy

Use Case	Best AWS Cache
Global CDN / HTTP caching	CloudFront
Reduce DB load	ElastiCache (Redis)
Faster DynamoDB reads	DAX
Web apps with static assets	CloudFront + S3
Real-time counters/leaderboards	Redis
Session store	Redis
Microsecond-latency lookups	Redis or DAX
Persisted in-memory data	Redis
Simple ephemeral caching	Memcached



Summary

AWS Caching Layers

1. **Edge caching** → CloudFront
2. **App/database caching** → ElastiCache (Redis/Memcached)
3. **DynamoDB acceleration** → DAX
4. **Static content caching** → S3 + CloudFront
5. **API-level caching** → API Gateway, CloudFront
6. **Serverless caching** → Lambda execution environment

1. Edge Caching with Amazon CloudFront

CloudFront is AWS's **global CDN**. It caches content **near users** at edge locations.

What CloudFront caches

- Static assets (HTML, CSS, JS, images)

- Dynamic content (with configurable TTL)
- API responses (with fine-grained cache keys)
- Lambda@Edge / CloudFront Functions transformed responses

Key CloudFront caching strategies

✓ Caching static content

- Set high TTL values
- Use versioned filenames (app.v123.js) to force refreshes

✓ Dynamic content caching

Configure:

- **Cache based on query strings**
- **Headers**
- **Cookies**
- **Authorization header bypass** (to avoid caching personal content)

✓ Origin Shield

Adds an extra caching layer to reduce origin load and improve cache hit rate.

✓ Signed URLs / cookies

To cache **private content** without exposing S3 origins publicly.

2. In-Memory Caching with Amazon ElastiCache (Redis / Memcached)

ElastiCache is used for **sub-millisecond access** and offloading heavy queries.

Caching strategies with ElastiCache

A. Cache-aside (Lazy loading) – Most common

Application checks cache → if miss → fetch from DB → store in cache.

Pros: Simple, effective

Cons: Cold-start on miss

B. Write-through

Write to cache **and** database simultaneously.

Pros: Cache always up-to-date

Cons: Writes slower, potentially unnecessary data stored

C. Write-behind (Write-back)

Application writes **only to cache**, and cache writes to DB asynchronously.

Pros: Very fast writes

Cons: Risk of data loss on failure; requires robust design

D. TTL-based caching

Heavy queries cached with expiration (e.g., 5-60 seconds).

Common use cases

- Session storage
- Leaderboards
- Real-time analytics
- Rate limiting
- Expensive join/query caching
- Token caching (OAuth, JWT blacklist)
- Distributed locks and queues (Redis)

Redis vs Memcached

Feature	Redis	Memcached
Data types	Rich (hashes, sets, lists, streams)	Key-value only
Persistence	Yes	No
Cluster mode	Yes	Yes
Use case	Complex caching, state, queues	Simple cache-only workloads

3. DynamoDB Accelerator (DAX)

DAX is a **fully managed Redis-like cache** for DynamoDB.

When to use it

- Read-heavy workloads
- Query patterns requiring microsecond latency
- You want a drop-in cache without code changes

Benefits

- Read latency goes from **ms** → **μs**
- No cache invalidation logic needed
- Auto-replicated cache cluster

4. Application + Database Caching Patterns

A. Query result caching

Store the result of expensive queries in ElastiCache or DAX.

B. Page/fragment caching

Cache fragments like menus, recommendations, etc.

C. Materialized view caching

Generate precomputed values and store in cache for fast retrieval.

D. Write-buffer / throttling via queues

Use SQS/Kinesis + cache to smooth spikes.

5. S3 + CloudFront for static content caching

A classic AWS pattern:

S3 (origin) → CloudFront (cache) → User

Best for:

- Static websites
- Media hosting
- Assets for mobile apps
- Video streaming

Configure:

- Cache-Control headers
- TTLs
- OAI / OAC for private S3 buckets

6. API and Microservices Caching

Using API Gateway + CloudFront

- API Gateway can cache responses at the API method level
- CloudFront adds global distribution and edge caching to APIs

Caching patterns:

- **Per-route caching**
- **Parameter-based cache keys**

- **Authorization-aware caching** (private APIs)

7. Lambda and Serverless Caching

Lambda Execution Environment Caching

Use global variables to persist between invocations:

- warm container → no re-fetch
- great for config, DB connections, small caches

Lambda + ElastiCache or DynamoDB DAX

For more robust, scalable caching.

3. How do you debug performance issues in an AWS environment?

High-Level Process

1. **Define the symptom** (slow requests? high latency? timeouts?)
2. **Locate the bottleneck** (compute, database, network, storage, cache, external services)
3. **Collect metrics + logs** using AWS tools
4. **Trace requests end-to-end**
5. **Validate resource configuration and scaling**
6. **Optimize or mitigate**

1. Use Key AWS Monitoring Tools

A. Amazon CloudWatch

Primary source for:

- CPU, memory, disk I/O
- Network throughput
- Lambda duration & cold starts
- RDS/ElastiCache metrics
- ALB/NLB request latency
- Autoscaling triggers
- Custom metrics

Focus on:

- **Latency (pXX)** – 90th, 95th, 99th percentiles
- **Throttling** (API Gateway, Lambda, DynamoDB)
- **SurgeQueue** (ALB under heavy load)
- **Read/write IOPS**
- **Database CPU / connections**

B. AWS X-Ray (Distributed Tracing)

Identifies:

- Slow microservices
- Slow DB queries
- High-latency external APIs
- Cold starts
- Event-driven delays (Lambda → SQS → Lambda)

X-Ray is your go-to for microservices debugging.

C. VPC Flow Logs (Networking Issues)

Useful for:

- Dropped traffic
- Latency due to NAT gateways
- Subnet routing misconfigurations
- Firewall (SG/NACL) issues

D. CloudTrail (Misconfigurations)

Detect:

- IAM changes
- Security group updates
- Network ACL modifications

2. Debug by AWS Layer

2.1 Compute: EC2 / ECS / EKS / Lambda

Common issues in compute services

- CPU or memory saturation
- Network throughput limits
- Insufficient instance types (IO-limited, CPU-limited)
- Lambda cold starts
- ECS/EKS pod resource constraints
- Auto Scaling too slow to react

Debugging steps

- Check CloudWatch CPU/Memory
- Inspect network metrics (packets out/in, ENI throttling)
- Look for throttling on Lambda concurrency
- Check container logs (CloudWatch Logs)

- Review ASG scaling events
- Use AWS Compute Optimizer for recommendations

2.2 Database: RDS / Aurora / DynamoDB

RDS/Aurora issues

- High CPU / low freeable memory
- Insufficient IOPS (storage bottleneck)
- Slow queries / missing indexes
- Locking / long-running transactions
- Connection saturation

Tools

- RDS Performance Insights
- Enhanced Monitoring
- Query Planner / EXPLAIN output

DynamoDB issues

- Hot partitions
- Provisioned throughput too low (RCU/WCU)
- Throttling (429 errors)
- Large item size

Tools

- DynamoDB Metrics (Throttles, Consumed Capacity)
- DAX metrics (if caching layer exists)

2.3 Networking: VPC / ALB / NLB / API Gateway

Symptoms

- Intermittent timeouts
- High tail-latency (p95/p99)
- Connection resets

Debugging

- Look at ALB TargetResponseTime
- Check SurgeQueueLength (ALB overloaded)
- Verify NAT Gateway throughput
- Inspect VPC Flow Logs
- Check API Gateway throttling or integration latency

2.4 Storage: S3, EFS, EBS

S3 Issues

- Slow prefix scanning
- High request rates hitting prefix limits
- Small object overhead

EBS/EFS Issues

- IOPS or throughput bottlenecks
- Bursting limits exhausted
- Wrong performance mode

Debugging

- CloudWatch metrics for IOPS, throughput
- fsx/S3 request metrics
- Enable S3 Server Access Logging

2.5 Caching Problems: CloudFront, ElastiCache, DAX

Typical issues

- Low cache hit ratio → backend overloaded
- Hot keys / skew
- TTL too low
- Evictions due to insufficient memory

Debugging

- CloudFront CacheHitRate
- Redis INFO (CPU, memory, evictions)
- DAX Miss/Hit ratios

2.6 CI/CD / Deployment Issues

Sometimes the issue is a configuration pushed via deployment:

- Incorrect autoscaling config
- Missing indexes in DB migrations
- Cold Lambda functions after massive rolling deploy
- New hot path in microservices

3. Identify the Bottleneck by Symptom

Symptom	Likely Cause
High latency	DB slow queries, throttling, cold starts, network issues

Symptom	Likely Cause
Timeouts	Load balancer overload, NAT gateway bottleneck, unscaled backend
Spikes in CPU	No autoscaling or burst workload
Intermittent failures	Connection pooling issues, DNS problems
Increased cost	Inefficient queries, misconfigured TTLs, noisy neighbors

4. Tools for Deep Debugging

AWS-native

- CloudWatch Logs Insights (log queries)
- CloudWatch Synthetics (canaries for testing endpoints)
- X-Ray Service Map
- Performance Insights (RDS)
- AWS CDK Diff / CloudFormation Drift Detection
- Compute Optimizer

Third-party

- Datadog
- NewRelic
- Splunk
- Prometheus/Grafana (EKS/ECS)

5. Common Fixes After Debugging

- Add caching (CloudFront, Redis, DAX)
- Add read replicas / faster DB instance
- Fix slow SQL queries (indexes, schema fixes)
- Adjust autoscaling thresholds
- Increase Lambda memory (improves CPU too)
- Use EBS io2 for consistent IOPS
- Reduce multi-hop network paths
- Re-architect monolith into smaller components

6. Golden Rule: Always Collect Data Before Changing Anything

Performance debugging is 80% **observability**, 20% **optimization**.

